

Scope0

Enterprise AI Agent Permission System — How it differs from CLAUDE.md

Scope0 vs CLAUDE.md File Permissions

CLAUDE.md files and Scope0 solve related but fundamentally different problems. This document explains where each fits, and why enterprise AI deployments require something beyond a markdown instruction file.

Capability	CLAUDE.md	Scope0
Enforcement	Soft — model tries to follow	Hard — blocked at API level
Identity & authN	None	JWT, API keys, SSO, MFA
Audit trail	None	Signed, tamper-evident log
Dynamic policy	Static text	Runtime evaluation per request
Multi-org / tenancy	Single session	Full org hierarchy + RBAC
Agent revocation	Not possible	Auto-revoke on anomaly
Compliance docs	None	Built-in audit & compliance API
Billing & quotas	None	Plan-based seat + decision limits

1. Hard Enforcement vs. Soft Guidance

A CLAUDE.md can say "don't touch production." Scope0 actually blocks the API call if the agent's token doesn't carry the required scope — the agent physically cannot proceed, regardless of the instructions it was given or how it interprets them.

2. Identity and Multi-Tenancy

CLAUDE.md has no concept of which agent is making a request, or which organisation owns it. Scope0 issues JWTs, manages org memberships, and enforces seat roles (Owner, Admin, Developer, Auditor, Member) so you know exactly who did what, for which customer, at what time.

3. Cryptographic Audit Trail

Every decision Scope0 makes — allow or deny — is recorded with a tamper-evident cryptographic chain. When a regulator asks whether an AI agent accessed customer data on a given day, you have a signed, verifiable answer. A CLAUDE.md leaves no record at all.

4. Dynamic Policy Engine

Scope0 evaluates context at request time: agent risk score, geographic origin, time of day, MFA status, active JIT grants. CLAUDE.md is static text authored before the session starts and cannot react to runtime signals.

5. Guardrails and Automatic Revocation

If an agent exhibits anomalous behaviour — rate spikes, unusual resource access, policy violations — Scope0's guardrail engine can auto-revoke the agent's permissions mid-session. CLAUDE.md has no runtime monitoring layer.

6. Enterprise Features: Billing, SSO, Compliance

Scope0 ships with plan-based seat and decision quotas, Okta/SAML SSO integration, role-restricted invite flows, and built-in compliance documentation APIs. These are requirements for selling into enterprise accounts — none of them are possible with a markdown file.

When CLAUDE.md Is the Right Tool

For a single developer using Claude Code on their own machine, CLAUDE.md is entirely sufficient — it's the right tool at that scope. It communicates intent, conventions, and project context to a single model session at zero infrastructure cost.

The Simple Analogy

CLAUDE.md	Scope0
Coding style guide	IAM / access control system
Tells Claude what to try	Decides what agents can do
Single session, one developer	Multi-org, thousands of agents
Best-effort compliance	Cryptographic enforcement